

# API Gateways & the Edge

Part II · Building Blocks — the single smart front door to your whole system.

## LEARNING OBJECTIVES

- Explain what an API gateway is and the problems it solves.
- List the gateway's key jobs: routing, authentication, rate limiting, aggregation, TLS, and more.
- Explain request aggregation and the Backend-for-Frontend (BFF) pattern.
- Distinguish an API gateway from a load balancer and a service mesh.
- Know the gateway's trade-offs and when a separate gateway isn't worth it.

## REAL-WORLD ANALOGY — THE BUILDING CONCIERGE


A large office building has one **concierge** desk that every visitor passes. The concierge checks who you are (authentication), stops anyone trying to flood the lobby (rate limiting), directs you to the right department (routing), and can even gather items from several offices so you don't walk the whole building yourself (aggregation). Visitors never wander the halls hunting for the right room. An **API gateway** is that concierge for all the traffic entering your system.

## 9.1 The problem the gateway solves

By now the system has grown: a load balancer, an app tier, caches — and soon many separate services (microservices, Part V). Without a single entry point, clients would have to know the address of every service, and each service would re-implement the same cross-cutting concerns: authentication, rate limiting, logging, TLS. That is duplicated, fragile, and couples clients to your internal structure. An **API gateway** is the single front door that centralizes all of this.

## 9.2 What an API gateway is

An API gateway is a server that sits at the **edge**, in front of all your backend services, and every client request goes through it. It is a specialized **reverse proxy** (Chapter 7) — but where a plain load balancer just distributes traffic, a gateway is **application-aware** and feature-rich, handling the cross-cutting concerns that every service would otherwise duplicate.



## 9.3 The key jobs of a gateway

| Job                    | What it does                                                                                                                        |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| <b>Routing</b>         | Sends each request to the correct backend service. Clients call one address; services can move or split without clients noticing.   |
| <b>Authentication</b>  | Validates the caller's token once, at the edge, rejecting bad requests before they reach any service (Chapters 5, 41).              |
| <b>Rate limiting</b>   | Caps how much any client can send, protecting backends from abuse and spikes (Chapter 35).                                          |
| <b>Aggregation</b>     | Turns one client call into several backend calls and merges the results — fewer round trips.                                        |
| <b>TLS termination</b> | Handles HTTPS at the edge so internal services can work in plain HTTP on a trusted network.                                         |
| <b>Cross-cutting</b>   | Central logging, metrics, caching, request/response transformation, and protocol translation (e.g. external REST to internal gRPC). |

### KEY INSIGHT — DO CROSS-CUTTING WORK ONCE, AT THE EDGE

Auth, rate limiting, TLS, and logging are needed by *every* service. Without a gateway you copy that logic into all of them and maintain it many times over. The gateway's core value is centralizing these concerns in one place, so each service can focus purely on its business logic.

## 9.4 Aggregation and Backend-for-Frontend

Mobile networks make round trips expensive (Chapter 3). If a screen needs data from four services, making the phone call all four is slow. With **aggregation**, the client makes *one* call to the gateway, which fans out to the four services in parallel and returns a single merged response.

This is often specialized as a **Backend-for-Frontend (BFF)**: a separate gateway tailored to each client type — one for mobile, one for web, one for TV — each shaping and combining exactly what that interface needs, instead of forcing one generic API to serve all of them.

## 9.5 Gateway vs load balancer vs service mesh

These three are easy to confuse; they solve different parts of the traffic problem.

| Component            | Role                                                                                                                            |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------|
| <b>Load balancer</b> | Distributes traffic across instances of a service (Chapter 7). Mostly “which server?”                                           |
| <b>API gateway</b>   | The smart <b>edge</b> front door for <i>incoming</i> client traffic (“north-south”): routing, auth, rate limiting, aggregation. |
| <b>Service mesh</b>  | Manages traffic <i>between internal services</i> (“east-west”): mTLS, retries, tracing — usually via sidecars (Chapter 30).     |

A gateway often *contains* a load balancer (it must still pick a healthy instance of each backend). The mesh handles what happens once traffic is already inside.

## 9.6 Trade-offs and pitfalls

### THE GATEWAY IS A CRITICAL COMPONENT

All traffic flows through it, so it’s a **single point of failure** and a potential **bottleneck** — always run several instances behind a load balancer, exactly as in Chapter 7. And resist stuffing business logic into it: a gateway that grows fat with domain rules becomes a “**distributed monolith**” front-end that every team must coordinate on, destroying the independence you built services to gain. Keep the gateway to cross-cutting concerns only.

## 9.7 When not to use a separate gateway

### DON’T ADD IT BEFORE YOU NEED IT

A single, simple backend often doesn’t need a dedicated gateway — your load balancer and web framework may already handle TLS, routing, and auth adequately. The gateway earns its place when you have **many services** or **many client types**, or when cross-cutting concerns are genuinely being duplicated. Like every block in this book, add it when the numbers and the pain justify it, not by default.

## Common mistakes

### WATCH OUT

- Running a **single** gateway instance — recreating a single point of failure.
- Putting **business logic** in the gateway, turning it into a distributed monolith.
- Letting each service re-implement auth and rate limiting instead of centralizing at the edge.
- Confusing the gateway (edge, north-south) with the service mesh (internal, east-west).
- Adding a heavy gateway to a simple single-service app that doesn’t need one.
- Forgetting the gateway adds a network hop — keep its own processing lean.

## Interview questions

1. What is an API gateway, and what problems does it solve?
2. List the key responsibilities of a gateway.
3. What is request aggregation, and what is the Backend-for-Frontend pattern?
4. How does a gateway differ from a load balancer and from a service mesh?
5. Why is the gateway a single point of failure, and how do you address that?
6. When would you *not* introduce a separate API gateway?
7. Why is doing authentication at the gateway better than in every service?

## Hands-on exercises

1. List which cross-cutting concerns you'd move to a gateway for a 5-service system.
2. Design one aggregated endpoint that a mobile home screen could call instead of four separate calls.
3. For each, say gateway or service mesh: validating a user's login token; encrypting service-to-service calls; routing `/api/v2` to a new backend.
4. Sketch a highly-available gateway deployment.
5. Decide whether a simple single-backend blog needs a separate gateway, and justify it.

## Mini project

### STAND UP A GATEWAY

Using a gateway (or Nginx configured as one), put a single entry point in front of two small services — say a “users” service and a “posts” service. Configure **path-based routing** (`/users/*` and `/posts/*`), add a simple **rate limit**, and terminate **TLS** at the gateway. Then add one **aggregated** endpoint that calls both services and merges the result. You'll have built the concierge from this chapter's hero diagram.

## Large project (capstone continues)

### YOUR SOCIAL PLATFORM — THE EDGE

Add an **API gateway** to your reference architecture as the single front door. Decide its responsibilities (routing to your services, token authentication, rate limiting, TLS termination), design one **aggregated** endpoint your mobile feed screen can call, and plan the gateway's own **high availability**. Note the boundary: the gateway handles incoming client traffic; internal service-to-service traffic will be handled by the mesh in Part V. Keep the gateway thin — cross-cutting concerns only.

## Chapter summary

- An **API gateway** is the single, application-aware front door at the edge — a feature-rich reverse proxy that every client request passes through.

- Its key jobs are **routing, authentication, rate limiting, aggregation, TLS termination**, and other cross-cutting concerns — done **once** so services don't each duplicate them.
- **Aggregation** and the **BFF** pattern cut client round trips by merging several backend calls into one.
- A gateway is the **edge** (north-south); a **load balancer** picks servers; a **service mesh** governs internal (east-west) traffic.
- Run **multiple** gateway instances (it's a critical path) and keep it **thin** — no business logic — and only add it when you genuinely need it.

#### COMING NEXT — CLOSING PART II

**Chapter 10 — Vertical vs Horizontal Scaling.** We've added load balancers, caches, and a gateway — now we formalize the two ways to grow, revisit the single-machine ceiling, and set up the hardest scaling problem of all: the data tier, which opens Part III.